

## Multiple Linear Regression and Geospatial

```
In [1]: # Check the dataset directory
        %pwd
```

```
Out[1]: 'C:\\Users\\SK\\Desktop\\Python\\Python\\Python Project\\5. Multiple Linear R
        egression and Geospatial'
```

```
In [2]: # Change the working directory
        import os
        os.chdir("/Users/SK/Desktop/SK/NUS EBA/Semester 2/Statistical BootCamp/WK4")
```

```
In [3]: # Import the functions
        import pandas as pd
        import numpy as np
        from pandas import DataFrame, read_csv
```

```
In [4]: # Read the csv file
        housing = pd.read_csv("housing.csv")
```

```
In [5]: housing.head()
```

```
Out[5]:
```

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households
0	-122.23	37.88	41.0	880.0	129.0	322.0	126.0
1	-122.22	37.86	21.0	7099.0	1106.0	2401.0	1138.0
2	-122.24	37.85	52.0	1467.0	190.0	496.0	177.0
3	-122.25	37.85	52.0	1274.0	235.0	558.0	219.0
4	-122.25	37.85	52.0	1627.0	280.0	565.0	259.0

In [6]: `housing.describe()`

Out[6]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	popul
<b>count</b>	20640.000000	20640.000000	20640.000000	20640.000000	20433.000000	20640.00
<b>mean</b>	-119.569704	35.631861	28.639486	2635.763081	537.870553	1425.47
<b>std</b>	2.003532	2.135952	12.585558	2181.615252	421.385070	1132.46
<b>min</b>	-124.350000	32.540000	1.000000	2.000000	1.000000	3.00
<b>25%</b>	-121.800000	33.930000	18.000000	1447.750000	296.000000	787.00
<b>50%</b>	-118.490000	34.260000	29.000000	2127.000000	435.000000	1166.00
<b>75%</b>	-118.010000	37.710000	37.000000	3148.000000	647.000000	1725.00
<b>max</b>	-114.310000	41.950000	52.000000	39320.000000	6445.000000	35682.00

In [7]: `housing.shape`

Out[7]: (20640, 10)

In [8]: *# Check the missing value*  
`housing.isnull().sum()`

Out[8]:

longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	207
population	0
households	0
median_income	0
median_house_value	0
ocean_proximity	0
dtype:	int64


In [13]: *# Remove the missing value from the dataset*  
`house_no_missing = housing.dropna()`

In [14]: `house_no_missing.isnull().sum()`

Out[14]:

longitude	0
latitude	0
housing_median_age	0
total_rooms	0
total_bedrooms	0
population	0
households	0
median_income	0
median_house_value	0
ocean_proximity	0
dtype:	int64

```
In [16]: import matplotlib.pyplot as plt  
house_no_missing.plot(kind="scatter", x="longitude", y="latitude", alpha=0.4)  
plt.show()
```



```
In [18]: house_no_missing.plot(kind="scatter", x="longitude", y="latitude",
    s=house_no_missing['population']/100, label="population",
    c="median_house_value", cmap=plt.get_cmap("jet"),
    colorbar=True, alpha=0.4, figsize=(10,7),
    )
plt.legend()
plt.show()
```

C:\Users\SK\Anaconda3\lib\site-packages\pandas\plotting\\_tools.py:307: MatplotlibDeprecationWarning:

The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get\_subplotspec().rowspan.start instead.

layout[ax.rowNum, ax.colNum] = ax.get\_visible()

C:\Users\SK\Anaconda3\lib\site-packages\pandas\plotting\\_tools.py:307: MatplotlibDeprecationWarning:

The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get\_subplotspec().colspan.start instead.

layout[ax.rowNum, ax.colNum] = ax.get\_visible()

C:\Users\SK\Anaconda3\lib\site-packages\pandas\plotting\\_tools.py:313: MatplotlibDeprecationWarning:

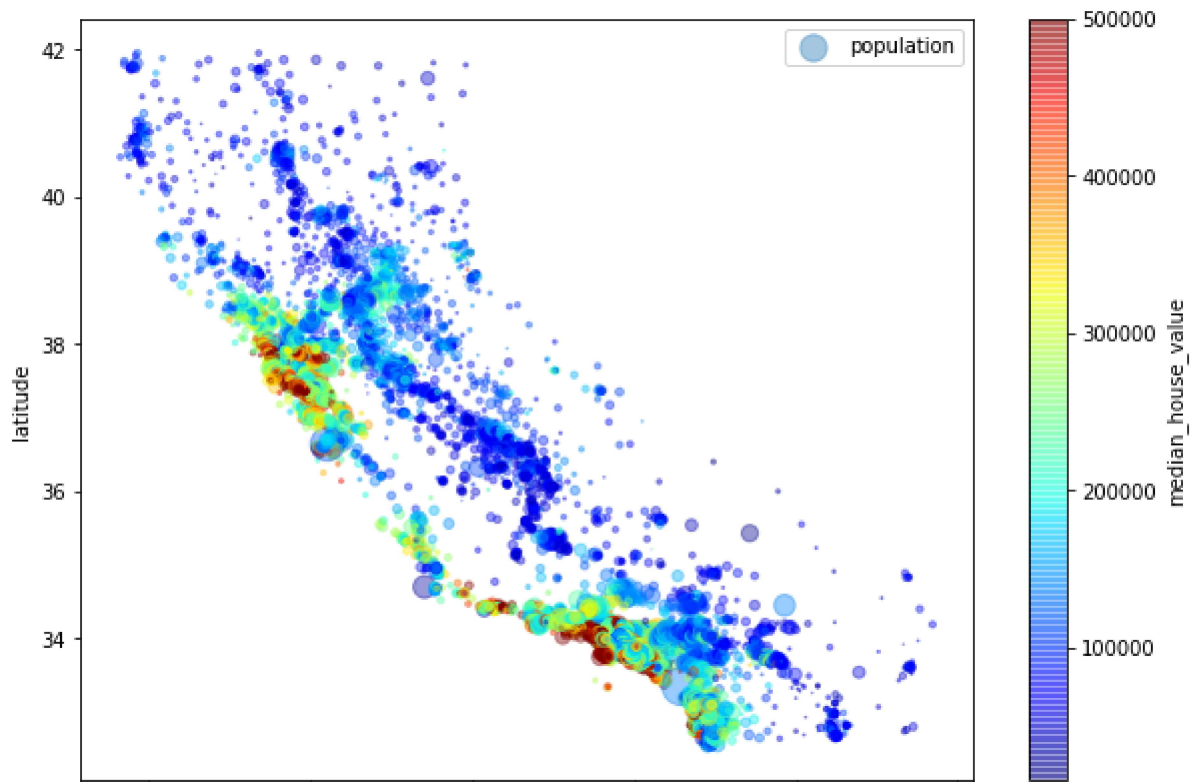
The rowNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get\_subplotspec().rowspan.start instead.

if not layout[ax.rowNum + 1, ax.colNum]:

C:\Users\SK\Anaconda3\lib\site-packages\pandas\plotting\\_tools.py:313: MatplotlibDeprecationWarning:

The colNum attribute was deprecated in Matplotlib 3.2 and will be removed two minor releases later. Use ax.get\_subplotspec().colspan.start instead.

if not layout[ax.rowNum + 1, ax.colNum]:



```
In [19]: from bokeh.io import output_file, output_notebook, show
from bokeh.models import (
    GMapPlot, GMapOptions, ColumnDataSource, Circle, LogColorMapper, BasicTic
    er,
    ColorBar, Range1d, PanTool, WheelZoomTool, BoxSelectTool)
from bokeh.models.mappers import ColorMapper, LinearColorMapper
from bokeh.palettes import Viridis5
```

```
In [20]: map_options = GMapOptions(lat=37.88, lng=-122.23, map_type="roadmap", zoom=6)

plot = GMapPlot(
    x_range=Range1d(), y_range=Range1d(),
    map_options=map_options
)
plot.title.text = "Hey look! It's a scatter plot on a map!"
```

```

In [21]: # For GMaps to function, Google requires you obtain and enable an API key:
#
#     https://developers.google.com/maps/documentation/javascript/get-api-key
#
# Replace the value below with your personal API key:
plot.api_key = "AIzaSyBYrbp340ohAHsX1cub8ZeHlMEFajv15fY"

source = ColumnDataSource(
    data=dict(
        lat=house_no_missing.latitude.tolist(),
        lon=house_no_missing.longitude.tolist(),
        size=house_no_missing.median_income.tolist(),
        color=house_no_missing.median_house_value.tolist()
    )
)
max_median_house_value = house_no_missing.loc[house_no_missing['median_house_v
alue'].idxmax()]['median_house_value']
min_median_house_value = house_no_missing.loc[house_no_missing['median_house_v
alue'].idxmin()]['median_house_value']

#color_mapper = CategoricalColorMapper(factors=['hi', 'lo'], palette=[RdBu3
[2], RdBu3[0]])
#color_mapper = LogColorMapper(palette="Viridis5", low=min_median_house_value,
high=max_median_house_value)
color_mapper = LinearColorMapper(palette=Viridis5)

circle = Circle(x="lon", y="lat", size="size", fill_color={'field': 'color',
'transform': color_mapper}, fill_alpha=0.5, line_color=None)
plot.add_glyph(source, circle)

color_bar = ColorBar(color_mapper=color_mapper, ticker=BasicTicker(),
                      label_standoff=12, border_line_color=None, location=(0,0
))
plot.add_layout(color_bar, 'right')

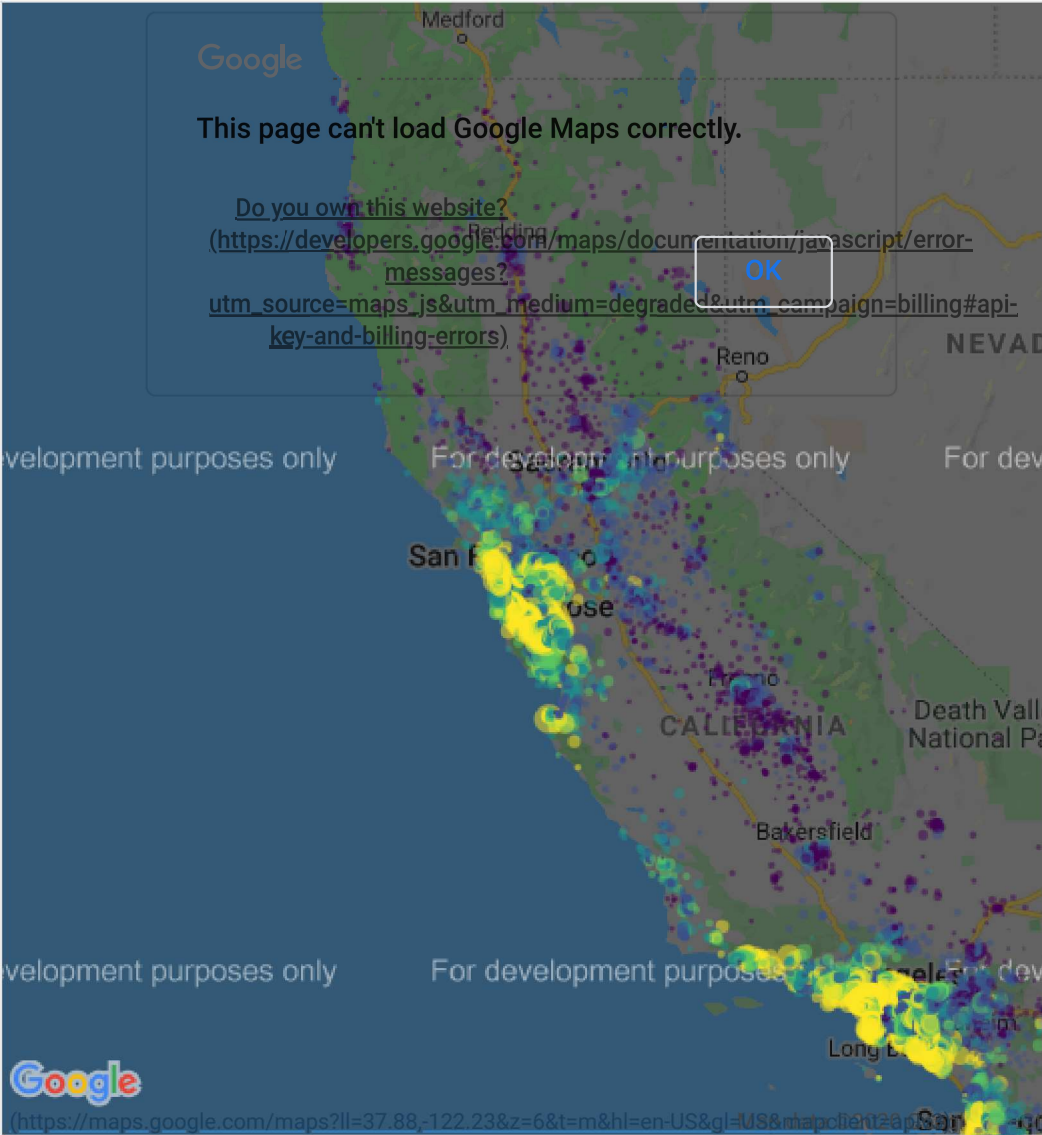
plot.add_tools(PanTool(), WheelZoomTool(), BoxSelectTool())
#output_file("gmap_plot.html")
output_notebook()

show(plot)

```

(https://www.kaggle.com/sk123/successfully-loaded)

Hey look! It's a scatter plot on a map!



(http

```
In [22]: ## Check the Correlation between variable
house_no_missing.corr()
```

Out[22]:

	longitude	latitude	housing_median_age	total_rooms	total_bedrooms	population	households	median_income	median_house_value
longitude	1.000000	-0.924616	-0.109357	0.045480	0.069608	0.100270	0.056513	-0.015550	-0.045398
latitude	-0.924616	1.000000	0.011899	-0.036667	-0.066983	-0.108997	-0.071774	-0.079626	-0.144638
housing_median_age	-0.109357	0.011899	1.000000	-0.360628	-0.320451	0.295787	0.302768	-0.118278	0.106432
total_rooms	0.045480	-0.036667	-0.360628	1.000000	0.930380	0.857281	0.918992	0.197882	0.133294
total_bedrooms	0.069608	-0.066983	-0.320451	0.930380	1.000000	0.877747	0.979728	-0.007723	0.049686
population	0.100270	-0.108997	0.295787	0.857281	0.877747	1.000000	0.979728	-0.007723	0.049686
households	0.056513	-0.071774	0.302768	0.918992	0.979728	0.979728	1.000000	-0.007723	0.049686
median_income	-0.015550	-0.079626	-0.118278	0.197882	-0.007723	-0.007723	-0.007723	1.000000	0.049686
median_house_value	-0.045398	-0.144638	0.106432	0.133294	0.049686	0.049686	0.049686	0.049686	1.000000

```
In [51]: ## Multiple Linear Regression
x = house_no_missing.drop(['median_house_value', 'ocean_proximity'], axis =1)
y = house_no_missing['median_house_value']
```



```
In [52]: from sklearn.linear_model import LinearRegression
import statsmodels.formula.api as smf
model = smf.ols('y~x', data = house_no_missing).fit()
model.summary()
```

Out[52]: OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.637
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.637
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	4478.
<b>Date:</b>	Wed, 15 Jul 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	23:46:43	<b>Log-Likelihood:</b>	-2.5682e+05
<b>No. Observations:</b>	20433	<b>AIC:</b>	5.137e+05
<b>Df Residuals:</b>	20424	<b>BIC:</b>	5.137e+05
<b>Df Model:</b>	8		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-3.585e+06	6.29e+04	-57.001	0.000	-3.71e+06	-3.46e+06
<b>x[0]</b>	-4.273e+04	717.087	-59.588	0.000	-4.41e+04	-4.13e+04
<b>x[1]</b>	-4.251e+04	676.952	-62.796	0.000	-4.38e+04	-4.12e+04
<b>x[2]</b>	1157.9003	43.389	26.687	0.000	1072.855	1242.945
<b>x[3]</b>	-8.2497	0.794	-10.387	0.000	-9.807	-6.693
<b>x[4]</b>	113.8207	6.931	16.423	0.000	100.236	127.405
<b>x[5]</b>	-38.3856	1.084	-35.407	0.000	-40.511	-36.261
<b>x[6]</b>	47.7014	7.547	6.321	0.000	32.909	62.493
<b>x[7]</b>	4.03e+04	337.207	119.504	0.000	3.96e+04	4.1e+04

<b>Omnibus:</b>	4898.534	<b>Durbin-Watson:</b>	0.975
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	18260.733
<b>Skew:</b>	1.166	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	7.002	<b>Cond. No.</b>	5.10e+05

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.1e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [55]: ## Check the VIF
from statsmodels.stats.outliers_influence import variance_inflation_factor
x['Intercept'] = 1

vif = pd.DataFrame()
vif["variables"] = x.columns
vif['VIF'] = [variance_inflation_factor(x.values, i) for i in range(x.shape[1])]

print(vif)
```

	variables	VIF
0	longitude	8.713740
1	latitude	8.828919
2	housing_median_age	1.260015
3	total_rooms	12.717000
4	total_bedrooms	36.003726
5	population	6.371238
6	households	35.136045
7	median_income	1.731511
8	Intercept	16702.386835

```
In [56]: ## Eliminate 'total_bedrooms' variable due to high VIF value
x = house_no_missing.drop(['median_house_value', 'ocean_proximity', 'total_bedrooms'], axis =1)
y = house_no_missing['median_house_value']
```

```
In [57]: model = smf.ols('y~x', data = house_no_missing).fit()
model.summary()
```

Out[57]: OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.632
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.632
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	5014.
<b>Date:</b>	Wed, 15 Jul 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	23:51:44	<b>Log-Likelihood:</b>	-2.5695e+05
<b>No. Observations:</b>	20433	<b>AIC:</b>	5.139e+05
<b>Df Residuals:</b>	20425	<b>BIC:</b>	5.140e+05
<b>Df Model:</b>	7		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-3.497e+06	6.31e+04	-55.434	0.000	-3.62e+06	-3.37e+06
<b>x[0]</b>	-4.197e+04	720.300	-58.273	0.000	-4.34e+04	-4.06e+04
<b>x[1]</b>	-4.222e+04	681.159	-61.983	0.000	-4.36e+04	-4.09e+04
<b>x[2]</b>	1126.4985	43.631	25.819	0.000	1040.979	1212.018
<b>x[3]</b>	-1.7751	0.694	-2.558	0.011	-3.135	-0.415
<b>x[4]</b>	-43.0960	1.052	-40.952	0.000	-45.159	-41.033
<b>x[5]</b>	148.9776	4.378	34.025	0.000	140.396	157.560
<b>x[6]</b>	3.838e+04	318.479	120.518	0.000	3.78e+04	3.9e+04

<b>Omnibus:</b>	5263.565	<b>Durbin-Watson:</b>	0.941
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	20453.770
<b>Skew:</b>	1.241	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	7.227	<b>Cond. No.</b>	5.01e+05

Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 5.01e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [58]: ## Check the VIF
x['Intercept'] = 1

vif = pd.DataFrame()
vif["variables"] = x.columns
vif['VIF'] = [variance_inflation_factor(x.values, i) for i in range(x.shape[1])]

print(vif)
```

	variables	VIF
0	longitude	8.677840
1	latitude	8.822928
2	housing_median_age	1.257568
3	total_rooms	9.583814
4	population	5.925295
5	households	11.673860
6	median_income	1.524467
7	Intercept	16579.355378

```
In [59]: ## Eliminate 'households' variable due to high VIF value
x = house_no_missing.drop(['median_house_value', 'ocean_proximity', 'total_bedrooms', 'households'], axis =1)
y = house_no_missing['median_house_value']
```

```
In [60]: model = smf.ols('y~x', data = house_no_missing).fit()
model.summary()
```

Out[60]: OLS Regression Results

<b>Dep. Variable:</b>	y	<b>R-squared:</b>	0.611
<b>Model:</b>	OLS	<b>Adj. R-squared:</b>	0.611
<b>Method:</b>	Least Squares	<b>F-statistic:</b>	5353.
<b>Date:</b>	Wed, 15 Jul 2020	<b>Prob (F-statistic):</b>	0.00
<b>Time:</b>	23:53:28	<b>Log-Likelihood:</b>	-2.5752e+05
<b>No. Observations:</b>	20433	<b>AIC:</b>	5.150e+05
<b>Df Residuals:</b>	20426	<b>BIC:</b>	5.151e+05
<b>Df Model:</b>	6		
<b>Covariance Type:</b>	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
<b>Intercept</b>	-3.968e+06	6.33e+04	-62.723	0.000	-4.09e+06	-3.84e+06
<b>x[0]</b>	-4.774e+04	719.613	-66.345	0.000	-4.92e+04	-4.63e+04
<b>x[1]</b>	-4.777e+04	679.805	-70.271	0.000	-4.91e+04	-4.64e+04
<b>x[2]</b>	1118.2926	44.848	24.935	0.000	1030.386	1206.199
<b>x[3]</b>	15.0384	0.501	30.020	0.000	14.057	16.020
<b>x[4]</b>	-25.4074	0.941	-27.014	0.000	-27.251	-23.564
<b>x[5]</b>	3.431e+04	303.318	113.100	0.000	3.37e+04	3.49e+04

<b>Omnibus:</b>	4604.156	<b>Durbin-Watson:</b>	0.816
<b>Prob(Omnibus):</b>	0.000	<b>Jarque-Bera (JB):</b>	12192.156
<b>Skew:</b>	1.215	<b>Prob(JB):</b>	0.00
<b>Kurtosis:</b>	5.902	<b>Cond. No.</b>	4.83e+05

#### Warnings:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 4.83e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
In [61]: ## Check the VIF
x['Intercept'] = 1

vif = pd.DataFrame()
vif["variables"] = x.columns
vif['VIF'] = [variance_inflation_factor(x.values, i) for i in range(x.shape[1])]

print(vif)
```

	variables	VIF
0	longitude	8.197084
1	latitude	8.316903
2	housing_median_age	1.257530
3	total_rooms	4.725402
4	population	4.479289
5	median_income	1.308668
6	Intercept	15780.636633